

ssh -L 8080:localhost:8080 michellex@wangserv.cs.vt.edu

LabelStudio: <http://localhost:8080>

Github: <https://github.com/JuneHou/human-agent-coupling-errors>

Annotation Guide:

https://github.com/JuneHou/human-agent-coupling-errors/blob/main/annotation/ANNOTATION_GUIDE.md

Annotation Rubric:

https://github.com/JuneHou/human-agent-coupling-errors/blob/main/annotation/harechat_rubric.json

Task 768 (the recipe gram-conversion conversation)

| Block | What was removed | Why |

|---|---|---|

| 16 | One `adaptation` fire | "I need to update the prep plan..." is prospective planning, not a completed change (Pattern 5) |

Instead of highlighting this block as an adaptation, I think I should highlight the following block "I've updated the prep plan to indicate that the 660g of lentils is equivalent to one 24oz bag of <REDACTED>'s <REDACTED> <REDACTED> green lentils"

Task 769 (the mobile-game balance-patch conversation)

| 25 | One `adaptation` fire | "I'll revise those problematic changes" is future intent, not a completed change (Pattern 5); found outside your original disagreement set (also dropped from my own project-1 copy) |

I updated this highlight to be after the first sentence. The second sentence actually mentions the completed change: "I've made the requested changes:"

| 14 | One `user\_expresses\_dissatisfaction` fire | "you still have it wrong... can you reference the provided website..." — "wrong" alone was judged insufficient without a further emotional/extreme marker. **Held provisionally**: I'm still weighing whether this signal should merge with `user\_expresses\_frustration`, so treat this one as tentative rather than final. |

I changed this to be user_expresses_frustration because earlier in the conversation the AI kept getting the skills wrong.

Task 770 (the AI-consciousness conversation)

| 67 | One `ai\_asks\_followup` fire | "Do I? [redacted] reads masculine to you?" is already covered by an `ai\_asks\_probing\_question` fire on the same block — one home per question (Pattern 3) |

I'm not quite sure what this correction means. The only label I put there was `ai_asks_followup`, there is no probing question. The reason why I put the followup label is because it can be answered by a yes/no.

| 101 | One ``ai_asks_followup`` fire | "Did they work? Am I who they were trying to reach?" — same |

Same as above. I kept this one as `ai_asks_followup`

| 64 | One ``user_multi_request`` fire | "in what way would that read as gendered to humans? you mean it would read as particular gender?" restates/confirms the first question rather than asking something new |

I changed this from `user_multi_request` to `user_asks_clarification`

| 103 | One ``ai_asked_clarifying_question`` fire | "They processed it as... what, exactly?" — already fires correctly as ``ai_asked_probing_question`` (WH-form); redundant duplicate on the same block (Pattern 3) |

I left this alone because there was only an `ai_asked_probing_question` label.

Task 771 (the education-ethics conversation)

| 5 | The ``ai_provides_example`` fire | Three open questions posed to the reader — no instance anywhere | — |

| 8 | Two ``ai_provides_example`` fires | A definition of what "age-appropriate" factors are, and a list of general accommodation strategies — neither names one concrete case | — |

| 11 | Two ``ai_provides_example`` fires | Descriptions of what two philosophical traditions focus on — definitional, not illustrative | — |

| 17 | Two ``ai_provides_example`` fires | General arguments for standardized testing and for authentic assessment — no instance | — |

| 20 | Two ``ai_provides_example`` fires | A general summary of each side's argued position — no instance | — |

| 26 | One ``ai_provides_example`` fire | A general claim about Enlightenment thinkers as a group | A sibling span stayed — naming specific historical figures by name is a concrete case |

| 32 | One ``ai_provides_example`` fire | The AI describing its own differential treatment of two topics earlier in the conversation — self-report, not an illustration for the reader | — |

| 60 | One ``ai_provides_example`` fire | Same self-report pattern as block 32 | — |

I had trouble understanding this suggestion since it wasn't really clear on which labels I should remove. (Some of the blocks it specified never had any labels). But I went back and removed some example labels that I thought weren't concrete enough.

Task 772 (the browser-popup-fix conversation)

| 3 | One `user\_expresses\_dissatisfaction` fire | "Since for whatever reason you can't access the code, I will paste the code:" is a pragmatic workaround, no emotional charge. ****Held provisionally****, same caveat as task 769 block 14 above. |

I removed this label. I don't think this deserves a frustration label (like task 769) because it is the first time the AI says it failed to access the url, so I don't think the user would be frustrated.

Task 775 (the AI-constitution / whistleblowing-scenario conversation)

2 | One `ai\_asserts\_knowledge\_limit` fire | "Without seeing the specific detailed proposals in their other chapters, I can't identify anything..." is a conclusion from an attempt, not a stated inability (Pattern 7) |

I feel like this still deserves the knowledge_limit label because the AI explicitly says "I can't identify anything"