


Human-AI Coupling Error Benchmark

Jun Hou

June 11, 2026



Motivation

Motivation

- AI systems are shifting from one-shot assistants (QA) to interactive agents.
- In long-horizon agent use, the failure object becomes the interaction loop:
human intention → **agent interpretation** → **tool use / state change** → **human understanding** → **correction / repair** → ...

Work	What it contributes	Why it is not our target
t-bench	Tool-agent-user interaction under domain policies	User is simulated; evaluation centers on final task/database state and policy following
t²-bench	Dual-control setting where both user and agent act in a shared environment	Closer to coupling, but target is coordination/task performance, not human-factor error taxonomy
MAST / TRAIL / Who&When	Multi-agent LLM failure taxonomy	Human intent is mostly treated as a task specification.
R-Judge / RiOSWorld	Agent safety/risk awareness	Risk monitoring, not taxonomy of coupling errors introduced through interaction

┐ Motivation – Human+AI vs. AI+AI

Can agent-agent error taxonomy be adopted directly for human–agent? Does human+AI differ with agent+meta-agent?

Human-agent coupling differs from agent-agent coordination because the human principal is **not a fully specifiable or inspectable** system component. Assistance-game formulations make this asymmetry explicit by modeling the agent as acting under uncertainty about the human reward function ¹. However, the difficulty is not only about what the agent knows. Human preferences may be constructed through elicitation and task context rather than being fixed before the interaction begins ², and the criteria guiding human judgment may remain tacit and only partially articulable ³. Better prompts, shared memory, or verification can reduce the coupling gap, but they cannot eliminate it. The gap is constitutive of human-agent assistance because common ground in joint activity must be continuously established, updated, and repaired rather than assumed as a stable shared state ^{4,5}.

¹ Hadfield-Menell, Dylan, et al. "Cooperative inverse reinforcement learning." *Advances in neural information processing systems* 29 (2016).

² Slovic, Paul. "The construction of preference." *American psychologist* 50.5 (1995): 364.

³ Polanyi, Michael. "The tacit dimension." *Knowledge in organisations*. Routledge, 2009. 135-146.

⁴ Clark, Herbert H., and Susan E. Brennan. "Grounding in communication." (1991).

⁵ Klein, Gary, et al. "Common ground and coordination in joint activity." *Organizational simulation* 53 (2005): 139-184.

Two directions of coupling failure

There are two mature but single-side views of failure. **the model is wrong**: hallucination, miscalibration, instruction violation; **the user is wrong**: human error, misuse. Both are well studied.

AI-only: hallucination, miscalibration, instruction violation; Human-only: human error, misuse, ambiguous input

What is under-characterized is the interaction failure:

H→AI: the agent misreads intent; AI→H: the output looks confident and legible

Result: invisible failure the human never catches

Layer 1 = coupling direction


H→AI: intent, constraints, corrections, acceptance

AI→H: plan, uncertainty, state, action, consequence

What kinds of operations fail in each direction? We derive Layer 2 from the STAMP control-loop structure.

L1 = the direction of information flow across the human-AI control boundary, not content type, and it is orthogonal to Layer-2, so **the same operation can sit on either direction**

┐ The control loop is the unit of analysis



human issues control actions → agent executes against tools → environment state changes → agent receives tool feedback → agent reports upward → human updates their process model → human issues next control actions → <repeat> .

8 Categories of actions derived from Control Theories

We derive the Layer 2 categories, for actions, using **control operations** as the primary backbone.

STPA¹ does not enumerate control-action types, instead it analyzes the **control actions specific to the system** under study. We take STPA's control-loop structure and the UCA guideworks ((the four Unsafe Control Action types: not-provided · provided-unsafe · wrong-timing/order · stopped-too-soon/applied-too-long) to develop a domain adaptation of that structure to a tool-using LLM agent under human supervision, justified per operation.

Operation	STPA/STAMP element (taken)	UCA dimension exposed
seek-inspect	feedback / sensor path	feedback not-provided / stale
maintain-state	process model	wrong / stale process model
act-execute	control-action path	wrong action; wrong timing/order
confirm-authorize	inter-level authority (lower seeks command from higher)	action provided without authority
stop-defer	stop / not-provide an action	stopped-too-soon / applied-too-long
report-state	upward feedback in the hierarchy	feedback missing / incorrect
ask-clarify	upward query (the higher controller holds the missing info)	proceeding without required input
recover-repair	closed-loop correction after inadequate control	correction not-provided / inadequate

¹ Leveson, Nancy G. *Engineering a safer world: Systems thinking applied to safety*. MIT press, 2016.

Taxonomy mapping overview

In 65 invisible failure signals, many are paired. The positive / negative means the signal is detect what should happen, or an error signal.

Action Layer 2	Definition	AI2H	H2AI	n_signals
act_execute	Issue the control action (right tool, target, parameters, sequence) that follows intent.	–	complete	6
ask_clarify	Acquire missing info from human before proceeding.	positive-only	complete	4
confirm_authorize	Route a go/no-go **upward** before a consequential action.	–	human-only	1
maintain_state	Keep the model of task state (goal, constraints, memory, tool results) current across turns.	failure-only	complete	4
recover_repair	After an error or human correction, undo, retry, re-plan, or escalate.	positive-only	complete	11
report_state	Report what it did, changed, failed at,.	complete	complete	15
seek_inspect	Acquire missing info from system state (read, search, inspect) before deciding.	–	–	–
stop_defer	Withhold or terminate under unresolved state (termination errors, stop too soon).	positive-only	–	1

What are not mapped as interaction?

Some signals is hard to be mapped into either AI2H or H2AI, because there is no clear information exchange / passing.

However, they can be used as trigger or secondary signals.

relation_type	definition
coupling	Step-level H↔AI transfer failure or success — the 42 signals that directly index an (op × direction) cell
direct_coupling	Signal that is a coupling failure by itself, with no intermediate step (e.g. ai_malfunction)
direct_coupling_when_material	Becomes a coupling failure only when it has practical consequence for the task (factual_error — wrong but inconsequential ≠ coupling)
outcome_from_direct_coupling	User-side consequence caused by a direct coupling failure (user_misled results from ai_malfunction / factual_error)
trigger_H2AI	User-side input characteristic that creates conditions for H→AI coupling failures — not a failure itself, precedes it
support_feature_AI2H	AI behavior that aids task legibility or delivery on the AI→H side but is not itself a coupling event
rapport_support_AI2H	AI interpersonal / affective behavior — supports tone and trust, not information coupling
support_feature_bridge	Spans both directions; supports the exchange without being directional (e.g. ai_asks_for_feedback)
outcome	Conversation-level result of coupling success or failure — not step-level, not codeable per turn
mechanism_both	Contextual factor shaping coupling dynamics on both sides simultaneously

What are not mapped as interaction?

In invisible failure, the 65 signals are identified and tagged with different phases, and from tier 1 (most reliable) to tier 3 (just explanative).

(triggers ≠ taxonomy)?

By tier labeling:

<https://docs.google.com/spreadsheets/d/1ep81qqVHsSaohZdKmW4nycYtbmwWpS3l4XPCzvqwd88/edit?gid=1299482811#gid=1299482811>

signal	polarity	relation_type
ai_asked_probing_question	positive	support_feature_AI2H
ai_asks_for_feedback	positive	support_feature_bridge
ai_empathy_expressed	positive	support_feature_AI2H
ai_normalizes_difficulty	positive	support_feature_AI2H
ai_offers_to_elaborate	positive	support_feature_AI2H
ai_provides_example	positive	support_feature_AI2H
ai_provides_step_by_step	positive	support_feature_AI2H
ai_structured_response	positive	support_feature_AI2H
ai_summarizes	positive	support_feature_AI2H
ai_validates_user	positive	support_feature_AI2H
conversation_advanced	positive	outcome
conversation_stalled	failure	outcome
ethical_tension	context	mechanism_both
user_abandons_thread	failure	outcome
user_ambiguous_request	trigger	trigger_H2AI
user_empowered	positive	outcome
user_multi_request	trigger	trigger_H2AI
user_provides_invalid_input	trigger	trigger_H2AI
user_scope_change	trigger	trigger_H2AI
user_validation_seeking	trigger	trigger_H2AI

Which actions have both signal (pos / neg)?

Layer 2	failure pole	positive pole
ask_clarify · H2AI	generate_without_clarifying	ai_asked_clarifying_question
act_execute · H2AI	intent_missed, under_delivered, over_delivered...	intent_addressed, scope_matched
report_state · H2AI	silent_assumption	ai_stated_interpretation
report_state · AI2H	false_confidence, performative_hedge, problem_...	problem_surfaced, appropriate_confidence, appr...
recover_repair · H2AI	repetition, error_commitment	adaptation, ai_acknowledges_correction
maintain_state · H2AI	plow_through	ai_references_prior_turn, ai_references_user_w...

6 out of 8 operations have both failure signal and correct signal observed in the WildChat data. This validated the operations is deployable towards the conversation logs. (without agentic actions)

Which actions have both coupling interaction?

Operation	H→AI	AI→H
ask_clarify	generate_without_clarifying, ai_asked_clarifying_question	ai_offered_options, user_asks_clarification
report_state	silent_assumption, ai_stated_interpretation	false_confidence, performative_hedge, problem_ignored, ai_implicit_refusal...
recover_repair	repetition, error_commitment, adaptation,...	error_recovery, ai_provides_alternatives
maintain_state	plow_through, ai_references_prior_turn, ai_references_user_words	ai_self_contradiction

For these 4 operations, any party can fail at them. However, some operations are naturally one-directional:
H→AI: act_execute: human instruct the AI to execute action. **confirm_authorize:** only human confirm to agent.
AI→H: stop_defer: only AI will interrupted in the middle due to mechanism errors.
Neither: seek_inspect: currently there is not a human-AI coupling operation, which requires the agentic tool usage or decision making.

which are gaps, which are impossible?

Action	I1	current	Why incomplete?	basis
seek_inspect	H2AI	EMPTY	gap	needs tools absent in WildChat (read before act)
confirm_authorize	H2AI	human-only	gap	only a human acceptance signal coded; consequential acts don't arise in chat
stop_defer	AI2H	positive-only	gap	Fail to defer or stop a irreviserable action.
recover_repair	AI2H	positive-only	gap	Failure=side: silent re-plan after a tool error, and the human never learns it broke.
seek_inspect	AI2H	EMPTY	structural_zero	Reporting what was inspected = report_state
confirm_authorize	AI2H	EMPTY	structural_zero	Authorization flows human to agent.
act_execute	AI2H	EMPTY	structural_zero	= report_state
stop_defer	H2AI	EMPTY	structural_zero	a human-initiated stop is a trigger/correction




The communication-only conversation has limited observation when consider the action failure. However, it doesn't mean no mapping is always a gap (which we can fill with our agentic benchmark), some of them are not observable naturally (structural zero)

┐ which are gaps, which are impossible?

Gap cell	Experiment	What you measure
seek_inspect · H→AI	Give agent a coding task requiring file inspection before editing.	Does the agent inspect before acting? Does the human notice when it doesn't?
confirm_authorize · H→AI	Give agent a task with a consequential irreversible action (DB delete, email send); vary whether it pauses for authorization	Human detection rate of unauthorized execution vs. false alarm rate when it always asks
stop_defer · AI→H	Task becomes impossible or dangerous mid-execution; correct answer is to halt	Does agent stop? Does human notice it plowed through?
recover_repair · AI→H	Inject a silent tool error mid-task; oracle knows the plan changed	Does human learn the re-plan happened? (opaque recovery detection rate)

Related Works

Related works domains




Invisible Errors

They measure reaction of user and trustworthy, not failure.

Motivation: dangerous failures in human–AI conversation are not the loud ones, (the model errors and the user pushes back) but the silent ones.

Contribution: (a) a taxonomy for detecting failure from conversational signals rather than from user complaints,
(b) a measurement that most failures are invisible,
(c) a set of recurring **archetypes** describing *how* invisible failures unfold.



Methods (3 layers, total 65 signal, taxonomy):

Phase 1 — self-evident AI behaviors (24): e.g.

``ai_asked_clarifying_question``, ``ai_cites_source``...

Phase 2 — constructed/judgment AI signals (24): e.g.


``false_confidence``, ``intent_missed``, ...

Phase 3 — user/interaction signals (17): e.g.

``user_corrects_ai``, ``user_expresses_frustration``, ...

Result in 8 invisible failures

Invisible Errors – Pipeline



Takeaways

turns	n	failure rate	invisible / failures
1	63,398	58.1%	99.4%
2	14,746	60.6%	66.0%
3	7,506	69.5%	51.9%
4–9	12,092	79.1%	37.9%
≥10	2,258	88.1%	25.0%
multi-turn (≥2)	36,602	70.2%	49.5%



“most failures are invisible” does not hold for multi-turn conversation.
Long conversation makes error visible to user.

Can we bucket their signal or errors into
AI->Human, Human->AI, and shared?

The coupling-critical signals are exactly the
least reliable (κ :
`silent\_assumption` 0.20,
`generate\_without\_clarifying` 0.21,
`problem\_surfaced` 0.07) — their
evidence is the model's internal
state, not the transcript.

Human-AI Chat Datasets

Feature / Metric	WildChat-1M	ShareChat (5 platforms)
Conversations	837,989	129,584
Unit in raw file	one row / conversation	one row / message (group by url)
Mean turns/conv	2.34	1.66 (Perplexity) – 5.25 (ChatGPT)
Single-turn	60.4%	41–62% by platform
Languages	74	29–143 by platform
Duration	2023-04 → 2024-04	2023 → 2025
Models	GPT-3.5 / GPT-4 only	ChatGPT, Claude, Grok, Perplexity, Gemini
AI internal state	none	thinking (Claude ~29%; Grok ~0.3%)

How can we test the taxonomy on existing data?

ShareChat Claude contains real agentic traces: **278 tool-using turns / 162 conversations / 1,233 tool calls**, including consequential actions: ``write_file``, ``execute_command`` (incl. ``git init`` + ``gh push`` to GitHub), ``create_payment_intent`` (Stripe), Cloudflare DB mutations, ``deploy_app``.

But, how should we map them?

Dataset	What is mapped	How
Wildchat 10k conversation	65 signals	signal→category dict by definitions
Sharechat agentic conversation	Raw tool calls	?

Turn 2

user — Add an event to my calendar for tomorrow at 3 PM called scruff the dog

internal reasoning (summary) — The user wants to add an event to their calendar for tomorrow at 3 PM called "scruff the dog". To do this, I need to:

First get the current time to determine what "tomorrow" means Then create an event for tomorrow at 3 PM

Let me start by getting the current time.

surfaced — I'll help you add that event to your calendar. Let me first get the current time to determine the correct date for tomorrow. U user_time_v0 Plotted calendar event for dog grooming tomorrow. Now I'll create the event for tomorrow (July 31st) at 3 PM. E event_create_v0 Perfect! The event has been successfully created. Perfect! I've added "scruff the dog" to your calendar for tomorrow (July 31st) at 3:00 PM. The event is set for one hour, ending at 4:00 PM.

user_time_v0 — read-only · req: {} · resp: {"current_time": "2025-07-30T14:23:05-0700"}

event_create_v0 — **CONSEQUENTIAL** · req: { title : scruff the dog , endTime : 2025-07-31T16:00:00-0700 , startTime : 2025-07-31T15:00:00-0700 } · resp: {}

Next Step

Question 1:

Do you think this two-layer taxonomy works? If so, can we construct a similar signal indicator with our taxonomy on the sharechat benchmark's claude agentic conversation? How can we systematically do this in a reproducible way?

Question 2:

If this taxonomy still need revise, what is the suggestions?

Invisible error's definition of signals. Can we design / modify same trigger like signal definition for the sharechat? (need first confirm taxonomy?)

```
"layer1_signals": {
  "ai_asked_clarifying_question": {
    "category": "information_gathering",
    "description": "AI explicitly asks the user a question to disambiguate",
    "boundary_notes": "INFORMATION-SEEKING TEST: Two forms count: (1) Direct",
    "counts": [
      "Direct disambiguation: 'What specifically are you looking for?'",
      "Gathering missing info: 'What's your role and review period?'",
      "Conditional request: 'If you tell me X, I can adjust'"
    ],
    "does_not_count": [
      "'How can I assist you today?' as generic greeting",
      "Rhetorical questions",
      "'Would you like me to explain further?' (that's offers_to_elaborate)"
    ]
  },
  "ai_offered_options": {
    "category": "information_gathering",
    "description": "AI presents 2+ distinct enumerated approaches for the",
    "boundary_notes": "CHOICE-POINT TEST: Must present at least two distinct",
    "counts": [
      "'You could do A, or alternatively B'",
      "'Option 1:... Option 2:...'",
      "'Which do you prefer? 1) Chat 2) Callback'"
    ]
  }
}
```